

CS 383 – Machine Learning

Support Vector Machines

Slides adapted from material created by E. Alpaydin
Prof. Mordohai, Prof. Greenstadt, Pattern Classification (2nd Ed.),
Pattern Recognition and Machine Learning

Objectives

- Support Vector Machines
 - Optimization Objective
 - Large Margin Intuition
 - Non-Linear SVMs

SVM Resources


- Burges tutorial
<http://research.microsoft.com/enus/um/people/cburges/papers/SVMTutorial.pdf>
- Shawe-Taylor and Christianini tutorial
<http://www.support-vector.net/icml-tutorial.pdf>
- Lib SVM
<http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- LibLinear
<http://www.csie.ntu.edu.tw/~cjlin/liblinear/>
- SVM Light
<http://svmlight.joachims.org/>
- Power Mean SVM
<https://sites.google.com/site/wujx2001/home/power-mean-svm>
- Matlab
 - `fitcsvm`
 - `predict`

SVMs

- Support Vector Machines (SVMs) are one of the most important developments in pattern recognition in recent years
- Elegant and successful


SVM Intuition

- Idea: Maximize distance to closest example (of each type)
 - For now we'll assume total separability




SVM Intuition

- Training data is just a subset of all possible data
 - Suppose hyperplane is close to sample X_i (open circle)
 - The *margin* is small
 - If we see a new sample close to X_i it may be on the wrong side of the hyperplane
- Therefore poor generalization




SVM Intuition

- Intuition: We want hyperplane as far as possible from any sample
- New samples close to old samples will then be classified correctly
- Good generalization




SVM – Linearly Separable Case

- Definition: Margin
 - The margin is twice the absolute value of distance b of the closest example to the hyperplane
- Our goal is to maximize the margin



Linear Discriminant Functions

- The equation of a linear hyperplane is
$$y = xw + w_0$$
- We can then create a **discriminant function** as
$$g(x) = xw + w_0$$
- We can then choose the class based on
 - $g(x) > 0 \rightarrow x$ in class 1
 - $g(x) \leq 0 \rightarrow x$ in class 2



Margin Formula

$$g(x) = xw + w_0$$

- So how do we find w, w_0 based on our training data in order to maximize the margin?
- As usual we'll minimize or maximize something
- We know we want to maximize the distance of the *closest samples* to our desired hyperplane
 - These closest samples are called the **support vectors**
 - Optimal hyperplane is completely defined by support vectors
 - Nice and compact representation!

Margin Formula

- Any sample *on* the hyperplane will have $g(x) = 0$
- So the distance (say the L1 distance) of any sample to the hyperplane is then:

$$d(x|W) = |g(x) - 0| = |xw + w_0|$$

- We can (somewhat arbitrarily) decide that the distance of any support vector to this hyperplane should be one:

$$\begin{aligned} |xw + w_0| &= 1 \quad \forall x \in \{\text{support vectors}\} \\ |xw + w_0| &> 1 \quad \forall x \notin \{\text{support vectors}\} \end{aligned}$$

Margin Formula

$$d(x|W) = |xw + w_0|$$

- So we need to turn this into a minimization or maximization problem.
- Let's divide $d(x|W)$ by the magnitude of w , $||w||$.


$$\frac{|xw + w_0|}{||w||}$$

- Then the closest samples (the support vectors) will have a value of

$$\frac{1}{||w||}$$

Margin Formula

- Two times this value $\frac{1}{||w||}$ gives us the **margin**.
- So we want to maximize $\frac{2}{||w||}$
- Subject to constraints
 - $xw + w_0 \geq 1$ for any positive example x
 - $xw + w_0 \leq -1$ for any negative example x
- We can convert this to a minimization problem
 - Minimize $J(w) = \frac{||w||^2}{2}$
 - Constrained to $Y_i(X_i w + w_0) \geq 1 \quad \forall (X_i, Y_i), Y_i \in \{-1, 1\}$
- $J(w)$ is a quadratic function, thus there is a single finite global minimum!



Quadratic Programming

- This is called *quadratic programming* problem
 - Minimize/maximize some quadratic function subject to linear constraints
- You can solve quadratic programming problems using an extended simplex method
 - Here's a resource:
 - <http://www.csse.monash.edu.au/~berndm/CSE460/Lectures/cse460-5b.pdf>
- When coding you will be allowed to use SVM packages in your code
 - Matlab: `fitcsvm`, `predict`
 - Many other languages: `Libsvm`
- But the next few slides will show you how to find solutions easily for a “simple” case

Finding Hyperplane

Minimize this

- To solve this constrained optimization problem we introduce *Lagrange multipliers* $a_n \geq 0$ for each constraint

- This gives us the equation:

$$L(w, a) = \frac{||w||^2}{2} - \sum_{n=1}^N a_n (Y_n (X_n w + w_0) - 1)$$

- The minus sign in front of the Lagrange multiplier is because while we're minimizing with respect to w , we're maximizing with respect to a

Distance from margin, which we want to maximize

Finding Hyperplane

$$L(w, a) = \frac{\|w\|^2}{2} - \sum_{n=1}^N a_n (Y_n (X_n w + w_0) - 1)$$

- As with all of our quadratic minimization/maximization problems we need to look at the derivatives
- Since we have two variables (a and w), let's take the derivative with respect to each (one at a time)
- Taking the derivative with respect to w (ignoring w_0 , we'll do this in a second..) we get
 - $\frac{\partial}{\partial w} L(w, a) = w - \sum_{n=1}^N a_n Y_n X_n$
 - Setting this to zero and solving for w ...

$$w = \sum_{n=1}^N a_n Y_n X_n$$

Finding Hyperplane

$$L(w, a) = \frac{\|w\|^2}{2} - \sum_{n=1}^N a_n (Y_n (X_n w + w_0) - 1)$$

- Taking it with respect to w_0 we get

- $\frac{\partial}{\partial w_0} L(w, a) = \sum_{n=1}^N a_n Y_n$

- Therefore

$$\sum_{n=1}^N a_n Y_n = 0$$

Finding Hyperplane

- Let's expand the original formula
- $L(w, a) = \frac{1}{2} ||w||^2 - \sum_{n=1}^N a_n (Y_n (X_n w + w_0) - 1)$
 - $= \frac{1}{2} ||w||^2 - \sum_{n=1}^N (a_n Y_n X_n w + Y_n w_0 - a_n)$
 - $= \frac{1}{2} ||w||^2 - (\sum_{n=1}^N a_n Y_n X_n w + \sum_{n=1}^N a_n Y_n w_0 - \sum_{n=1}^N a_n)$
- From our derivatives we know that at the minimum
 - $w = \sum_{n=1}^N a_n Y_n X_n$
 - $\sum_{n=1}^N a_n Y_n = 0$

Finding Hyperplane

- So let's substitute in $w = \sum_{n=1}^N a_n Y_n X_n$ into our original formula

$$L(w, a) = \frac{1}{2} \|w\|^2 - (\sum_{n=1}^N a_n Y_n X_n w + \sum_{n=1}^N a_n Y_n w_0 - \sum_{n=1}^N a_n)$$

- $L(a) = \frac{1}{2} (\sum_{n=1}^N a_n Y_n X_n)^T (\sum_{n=1}^N a_n Y_n X_n) -$
 $(\sum_{n=1}^N a_n Y_n X_n (\sum_{n=1}^N a_n Y_n X_n) + \sum_{n=1}^N a_n Y_n w_0 - \sum_{n=1}^N a_n)$
- $L(a) = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i^T X_j -$
 $(\sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i^T X_j + \sum_{n=1}^N a_n Y_n w_0 - \sum_{n=1}^N a_n)$
- $L(a) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i^T X_j + \sum_{n=1}^N a_n Y_n w_0$

Finding Hyperplane

$$L(a) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i^T X_j + \sum_{n=1}^N a_n Y_n w_0$$

- We can also use the fact that at the minimum $\sum_{n=1}^N a_n Y_n = 0$ to eliminate the last term
- This results in

$$L(a) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i^T X_j$$

- We have transformed this into a minimization problem of only a single variable!

Finding Hyperplane

- Finally we take the derivative of this equation with regards to each of the variables a_n
- Setting these to zero we get a system of N equations
- After solving, all vectors (observations) associated with non-zero Lagrange multipliers are support vectors
- Using the solved Lagrange multipliers we can substitute back to solve for w since:

$$w = \sum_{n=1}^N a_n Y_n X_n$$

And then solve for w_0 since $|xw + w_0| = 1$ for all support vectors

Simple Example

- Find the support vectors for a simple case.
- You'll need:

$$L(a) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i X_j^T$$
$$w = \sum_{n=1}^N a_n Y_n X_n$$

- Follow this example....

$$X = \begin{bmatrix} -1 & -1 \\ -1 & 1 \\ 1 & -1 \\ 1 & 1 \end{bmatrix} \quad Y = \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$

Example

$$X = \begin{bmatrix} -1 & -1 \\ -1 & 1 \\ 1 & -1 \\ 1 & 1 \end{bmatrix} \quad Y = \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$

- Writing the minimization problem out for $N = 4$ we have:
 - $L(a) = (a_1 + a_2 + a_3 + a_4) - \frac{1}{2}(2a_1^2 + 0a_1a_2 - 0a_1a_3 + 2a_1a_4 + 0a_2a_1 + 2a_2^2 + 2a_2a_3 + 0a_2a_4 + 0a_3a_1 + 2a_3a_2 + 2a_3^2 + 0a_3a_4 + 2a_4a_1 + 0a_4a_2 + 0a_4a_3 + 2a_4^2)$
- Collecting terms...
 - $L(a) = (a_1 + a_2 + a_3 + a_4) - \frac{1}{2}(2a_1^2 + 2a_1a_4 + 2a_2^2 + 2a_2a_3 + 2a_3a_2 + 2a_3^2 + 2a_4a_1 + 2a_4^2)$

$$L(a) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i X_j^T$$

Example

$$X = \begin{bmatrix} -1 & -1 \\ -1 & 1 \\ 1 & -1 \\ 1 & 1 \end{bmatrix} \quad Y = \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix}$$

$$L(a) = (a_1 + a_2 + a_3 + a_4) - \frac{1}{2}(2a_1^2 + 2a_1a_4 + 2a_2^2 + 2a_2a_3 + 2a_3a_2 + 2a_3^2 + 2a_4a_1 + 2a_4^2)$$

- Now let's take the derivative of this with regards to each variable
 - $\frac{d}{da_1} = 1 - \frac{1}{2}(4a_1 + 2a_4 + 2a_4) = 1 - 2a_1 - 2a_4 = 0$
 - $\frac{d}{da_2} = 1 - \frac{1}{2}(4a_2 + 2a_3 + 2a_3) = 1 - 2a_2 - 2a_3 = 0$
 - $\frac{d}{da_3} = 1 - \frac{1}{2}(4a_3 + 2a_2 + 2a_2) = 1 - 2a_3 - 2a_2 = 0$
 - $\frac{d}{da_4} = 1 - \frac{1}{2}(4a_4 + 2a_1 + 2a_1) = 1 - 2a_4 - 2a_1 = 0$
- Setting these equations equal to zero we get the following system of equations:
 - $2a_1 + 2a_4 = 1$ (twice)
 - $2a_2 + 2a_3 = 1$ (twice)

Example

$$2a_1 + 2a_4 = 1 \text{ (twice)}$$

$$2a_2 + 2a_3 = 1 \text{ (twice)}$$

- Since there's only 2 independent equations and 4 variables there are multiple solutions.
- However each should give us the same hyperplane

$$w = \sum_{n=1}^N a_n Y_n X_n$$


- $a_1 = 1/2, a_4 = 0, a_2 = 0, a_3 = 1/2$
 - Only the vectors associated with a_1 and a_3 are support vectors
 - $w = -1/2[-1, -1] + 1/2[1, -1] = [1, 0]$
 - $|X_1 w| = |[-1, -1][1, 0]^T| = 1$ so $w_0 = 0$
- $a_1 = a_2 = a_3 = a_4 = 1/4$
 - So all four vectors are support vectors
 - $w = -1/4[-1, -1] - 1/4[-1, 1] + 1/4[1, -1] + 1/4[1, 1] = [1, 0]$
 - $|X_1 w| = |[-1, -1][1, 0]^T| = 1$ so $w_0 = 0$

Example

- Sanity check!
- All support vectors should have $|xw + w_0| = 1$
- All non-support vectors should have $|xw + w_0| > 1$


SVM: Non-Separable Case

- Data is most likely to be not linearly separable
 - But linear classifier may still be appropriate
- There are basically two approaches to this issue:
 - Allow for some margin of error
 - Move the data to a *higher* dimensionality where it may be separable



Non-Separable Linear SVM

- First we'll look at keeping the data in its current dimensionality but add a margin for error
 - Still the data should be “almost” linearly separable for good performance
- The solution uses *slack variables* ξ_1, \dots, ξ_N (one for each sample)
 - Which is essentially the amount of error (how far from the margin) for each sample given the current plane



Non-Separable Linear SVM


- So we minimize

$$J(w, \xi) = \frac{\|w\|^2}{2} + \beta \sum_{i=1}^N \xi_i$$

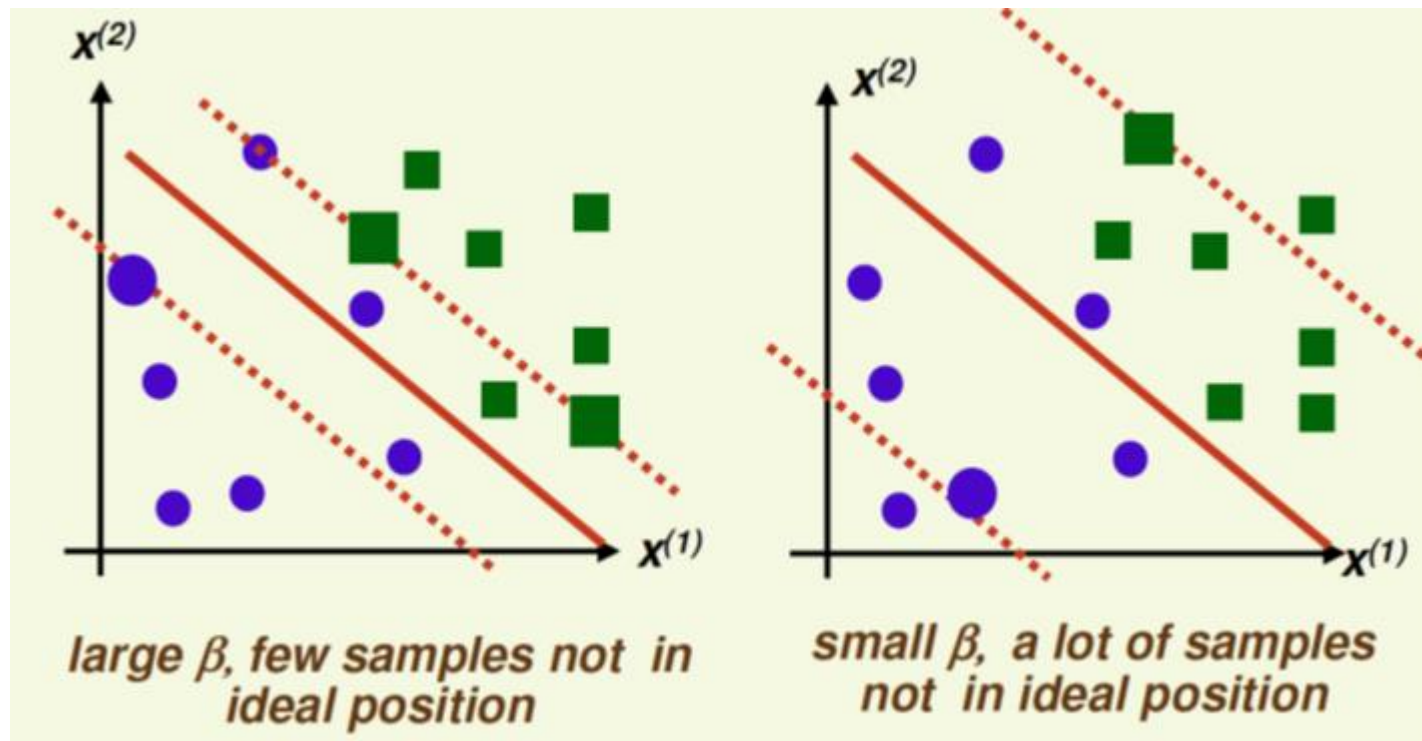
- Subject to:

- $\forall i, Y_i(X_i w + w_0) \geq 1 - \xi_i$
- $\forall i, \xi_i \geq 0$

- This can be interpreted as maximizing the width (although we inverted it to make it a minimization problem) while minimizing the amount of miss-classifications
- β is a constant that measure the relative weight of the first and second term
 - If β is small, we allow a lot of samples to be in not ideal positions (closer to linear SVM)
 - If β is large then the original non-ideal samples may be closer to the margin (better) but new ones may now be inside it.



SVM: Non-Separable Case



Finding Hyperplane


- If we add in our Lagrange multipliers, take the derivatives with respect to w , w_0 and ξ , and do substitution we get the equation:

$$L(a) = \sum_{n=1}^N a_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j X_i X_j^T$$

- Subject to
 - $0 \leq a_i \leq \beta$
 - $\sum_{i=1}^N a_i Y_i = 0$
- This is the same as our linearly separable case!
 - The only difference is the upper bound constraint on the values of a_i

Nonlinear SVM (kind of)

- SVMs define hyperplanes
 - So it's a linear classifier
- How can we deal with data that isn't linear separable?
 - We can try to deal with it in a higher dimension!
 - Doing so hopefully makes the data more separable



Non-Linear SVM

- Typically SVMs use the kernel trick to simulate computing the cosine similarity in higher dimensions without actually going to higher dimensions.
- Usually the choice is between a linear and an RBF kernel, the latter of which is non-linear
- As a rule of thumb
 - If the number of features is large (larger than number of observations) you may not need to map to a higher dimensional space therefore a linear kernel may be suffice
 - Plus it can become computationally expensive
 - Otherwise use RBF

Nonlinear SVM Step-by-Step

1. Start with data X_1, \dots, X_n in D -dimensional space
2. Choose kernel κ
 - From this you should know its mapping function $\phi(x)$
3. Choose a blending value β (if you want)
4. Find maximum margin linear discriminant function in higher dimensional space by solving quadratic programming problem (use some package)

$$L(a) = \sum_{n=1}^N a_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N a_i a_j Y_i Y_j \kappa(X_i, X_j)$$

- Constrained to $0 \leq a_i \leq \beta \forall i$ and $\sum_{i=1}^n a_i Y_i = 0$
- This will give us S is the set of support vectors

$$S = \{X_i | a_i \neq 0\}$$

Nonlinear SVM Step-by-Step

5. Given new vector x , determine which class it belongs to using the linear discriminant function


$$g(x) = \phi(x)w + w_0$$

- However, due to the kernel trick we can actually just compute this using the support vectors and the kernel function without ever computing $\phi(x)$ or finding w or w_0

$$g(x) = \sum_{X_i \in S} a_i Y_i \kappa(X_i, x)$$


- Note again we don't actually need to project x into higher dimensional space to do the computation
- $g(x) > 0$ class 1, otherwise class 2

SVM Example: XOR



- Class 1 (positive): $X_1=[1,-1]$, $X_2=[-1,1]$
- Class 2 (negative): $X_3=[1,1]$, $X_4=[-1,-1]$
- Kernel choice: Polynomial of degree 2
 - $\kappa(x, y) = (xy^T + 1)^2$
 - So $\phi(x) = [1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2, x_1^2, x_2^2]$
- Need to maximize
 - $L_D(a) = \sum_{i=1}^4 a_i - \frac{1}{2} \sum_{i=1}^4 \sum_{j=1}^4 a_i a_j Y_i Y_j (X_i X_j^T + 1)^2$
 - Constrained to $0 \leq a_i \forall i$ and $\sum Y_i \alpha_i = 0$
- Using a quadratic programming package we get:
 - $a_1 = a_2 = a_3 = a_4 = \frac{1}{8}$
 - All samples are support vectors since they are all non-zero

SVM Example: XOR



- Let's test our sample $[1,1]$
- Recall $\kappa(x, y) = (xy^T + 1)^2$
- $g(1,1) = \frac{1}{8}(+)\kappa([1, -1], [1,1]) + \frac{1}{8}(+)\kappa([-1, 1], [1,1]) + \frac{1}{8}(-)\kappa([1,1], [1,1]) + \frac{1}{8}(-)\kappa([-1, -1], [1,1])$
- $= \frac{1}{8}(0 + 1)^2 + \frac{1}{8}(0 + 1)^2 - \frac{1}{8}(3)^2 - \frac{1}{8}(-2 + 1)^2$
- $= \frac{1}{8} + \frac{1}{8} - \frac{9}{8} - \frac{1}{8} = -1$

SVM Example: XOR Problem

- If we actually were interested in the weight vector we could compute it (albeit with the help of ϕ):

$$\begin{aligned}w &= \sum_{i=1}^n a_i Y_i \phi(X_i) \\&= \frac{1}{8} (+) \phi(X_1) + \frac{1}{8} (+) \phi(X_2) + \frac{1}{8} (-) \phi(X_3) + \frac{1}{8} (-) \phi(X_4)\end{aligned}$$

- Recall that we're using

$$\phi(x) = [1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2, x_1^2, x_2^2]$$

- So with some arithmetic:

$$w = \left[0, 0, 0, -\frac{\sqrt{2}}{2}, 0, 0 \right]^T$$


SVM Example: XOR Problem

- Let's find w_0
 - $|\phi(X_s)w + w_0| = 1$ for any support vector X_s
 - Let's use support vector $X_s = [1, 1]$
 - $[1, 1, \sqrt{2}, \sqrt{2}, \sqrt{2}, 1, 1] \left[0, 0, 0, -\frac{\sqrt{2}}{2}, 0, 0\right]^T = -1$
 - So $w_0 = 0$
- Finally we can say
$$g(x) = \phi(x)w + w_0 = -x_1x_2$$


SVM Example: XOR Problem

Error: $g(x) = -x_1x_2$

$$g(x) = -2x^{(1)}x^{(2)}$$



decision boundaries nonlinear



decision boundary is linear

Matlab: `svmtrain/svmclassify`

- If you read the Matlab `fitcsvm` documentation you'll see that you can choose the kernel function type
 - And specify parameters for the chosen kernel
- In addition, after training your svm model contains:
 - `IsSupportVector` – A Boolean vector saying whether or not each observation is a support vector.
 - `Alpha` – Vector of weights for the support vectors
 - This is basically w
 - `Bias` – w_0
- You can then use `predict` to classify observations
 - This just gives you a class ID
 - If you want a distance, you need to compute it yourself using `Alpha` and `Bias`!

SVM Summary

- Pros
 - Strong theory
 - Good generalization
 - Global minima/maxima
- Cons
 - Directly applicable only to binary classification
 - Quadratic programming is expensive
 - Need to choose kernel

Final Observations

- Let's think about this algorithm
 - Supervised or non?
 - Classification or regression?
 - Model-based or instance-based?
 - When it comes time to test/use, are we using the original data?
 - Linear vs Non-Linear?
 - Can this work on categorical data?
 - Can this work on continuous valued data?
 - Training Complexity?
 - Testing Complexity?
 - How to deal with overfitting?
 - Directly handles multi-class?